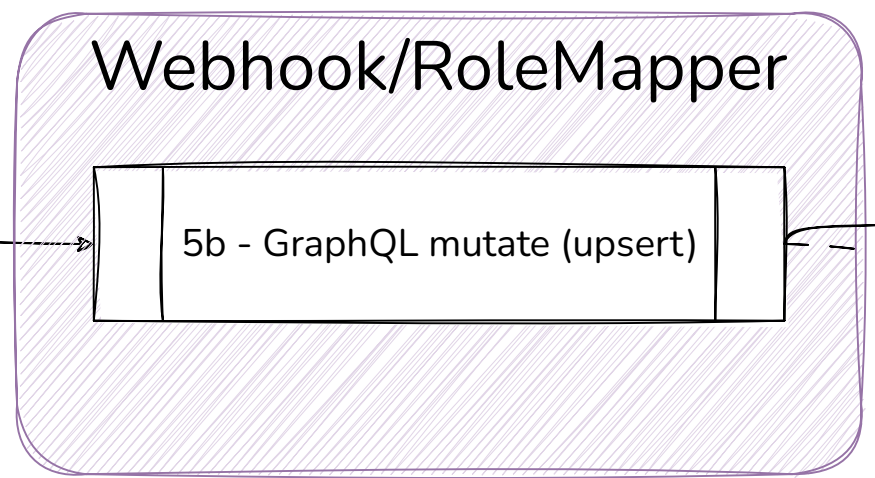
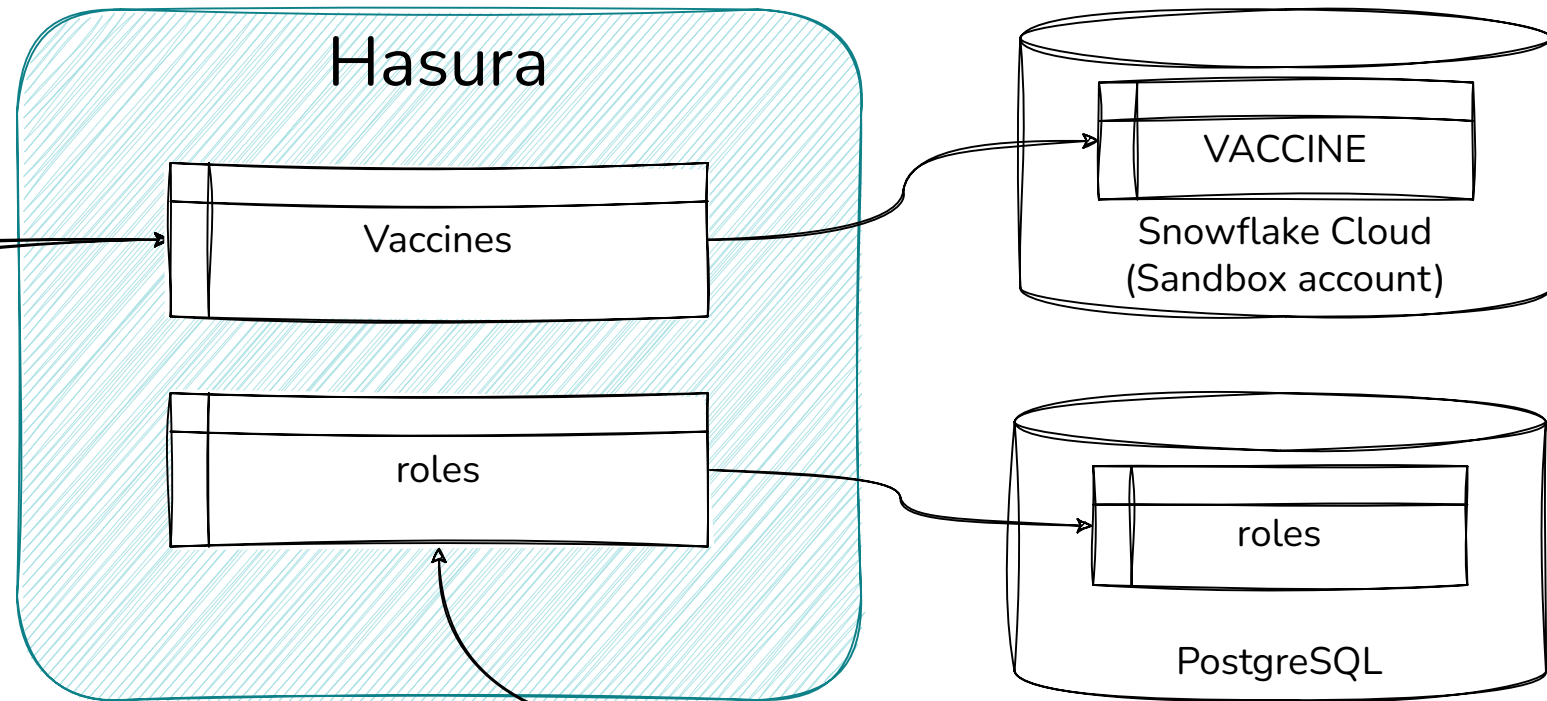
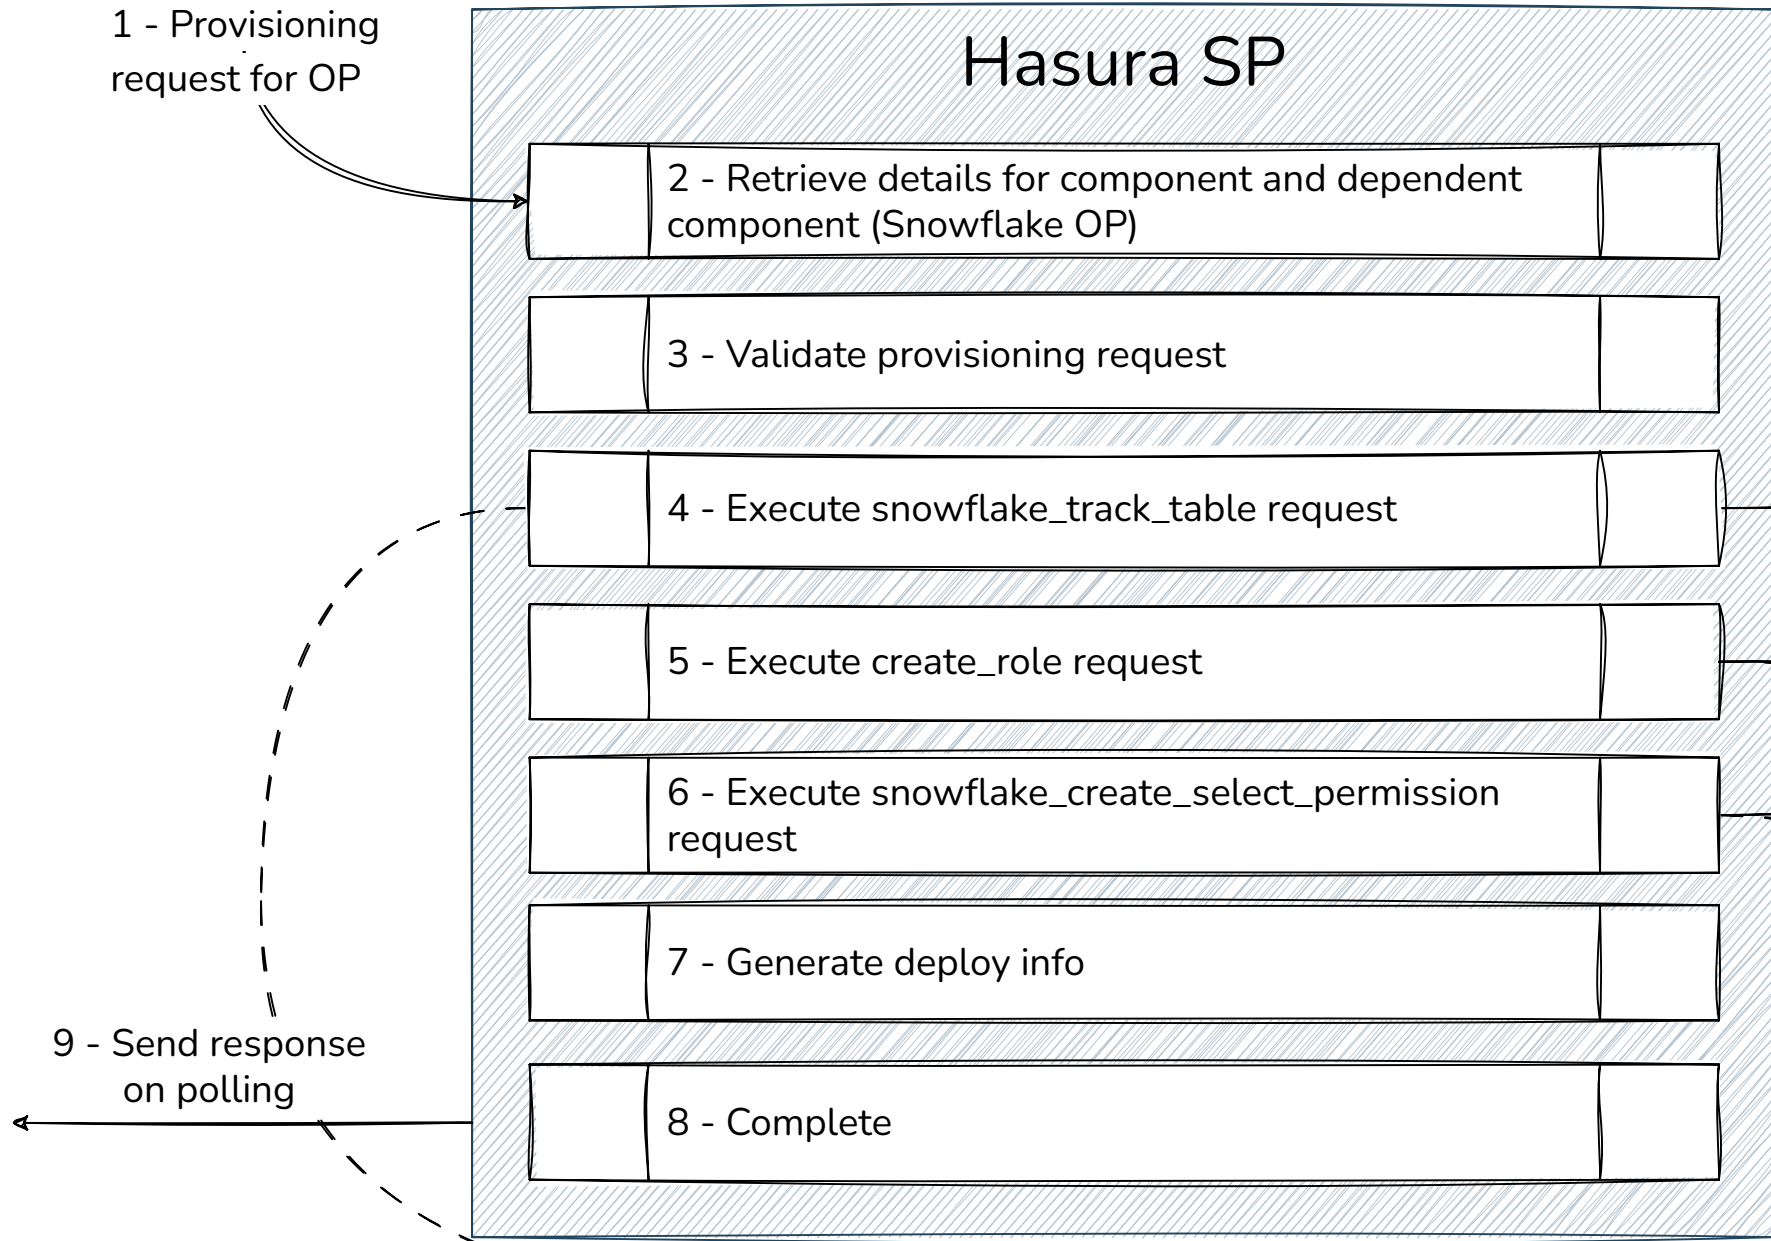


Idempotent, ie does not fail if something it provisions already exists.

The corresponding unprovisioning operation is very simple: just untrack the table.

- Step 4 changes to an untrack
- Step 5 becomes role delete
- Step 6 disappears

We assume the Data Sources exist already



**Validations:**

- Check data source exists on Hasura
- Check Snowflake OP is from the same DP
- Check configuration for root fields is valid for GraphQL

```
POST v1/metadata
Content-Type: application/json
X-Hasura-Role: admin
X-Hasura-Admin-Secret: myadminsecretkey
{
  "type": "snowflake_track_table",
  "args": {
    "source": "snowflake",
    "table": "VACCINE",
    "configuration": {
      "custom_name": "Vaccines",
      "custom_root_fields": {
        "select": "Vaccines",
        "select_by_pk": "Vaccine",
        "select_aggregate": "VaccineAgg",
        "insert": "AddVaccines",
        "insert_one": "AddVaccine",
        "update": "UpdateVaccines",
        "update_by_pk": "UpdateVaccine",
        "delete": "DeleteVaccines",
        "delete_by_pk": "DeleteVaccine"
      },
      "comment": "Access to VACCINE table"
    },
    "apollo_federation_config": {
      "enable": "v1"
    }
  }
}
```

```
POST v1/metadata
Content-Type: application/json
X-Hasura-Role: admin
X-Hasura-Admin-Secret: myadminsecretkey
{
  "type": "snowflake_create_select_permission",
  "args": {
    "table": [
      "VACCINE"
    ],
    "role": "dom1.dp1.0.op.readrole",
    "permission": {
      "columns": [
        ],
      "filter": {
        ],
      "set": [
        ],
      "allow_aggregations": false
    },
    "source": "snowflake"
  }
}
```

```
POST v1/role
Content-Type: application/json
{
  "role_id": "dom1.dp1.0.op.readrole",
  "component_id": "urn:dmb:cmp:dom1:dp1:0:op",
  "graphql_root_field_name": "dom1_dp1_0_op"
}
```

```
POST v1/graphql
Content-Type: application/json
X-Hasura-Role: admin
X-Hasura-Admin-Secret: myadminsecretket
{
  "query": "mutation InsertRole($component_id: String = \"\", $graphql_root_field_name: String = \"\", $role_id: String = \"\") {\n  insert_roles_one(object: {component_id: $component_id, graphql_root_field_name: $graphql_root_field_name, role_id: $role_id}, on_conflict: {constraint: roles_pkey, update_columns: [component_id, graphql_root_field_name], where: {graphql_root_field_name: {_eq: $graphql_root_field_name}, _and: {component_id: {_eq: $component_id}}}) {\n  component_id\n  graphql_root_field_name\n  role_id\n  }\n}",
  "variables": {
    "role_id": "dom1.dp1.0.op.readrole",
    "component_id": "urn:dmb:cmp:dom1:dp1:0:op",
    "graphql_root_field_name": "dom1_dp1_0_op"
  },
  "operationName": "InsertRole"
}
```

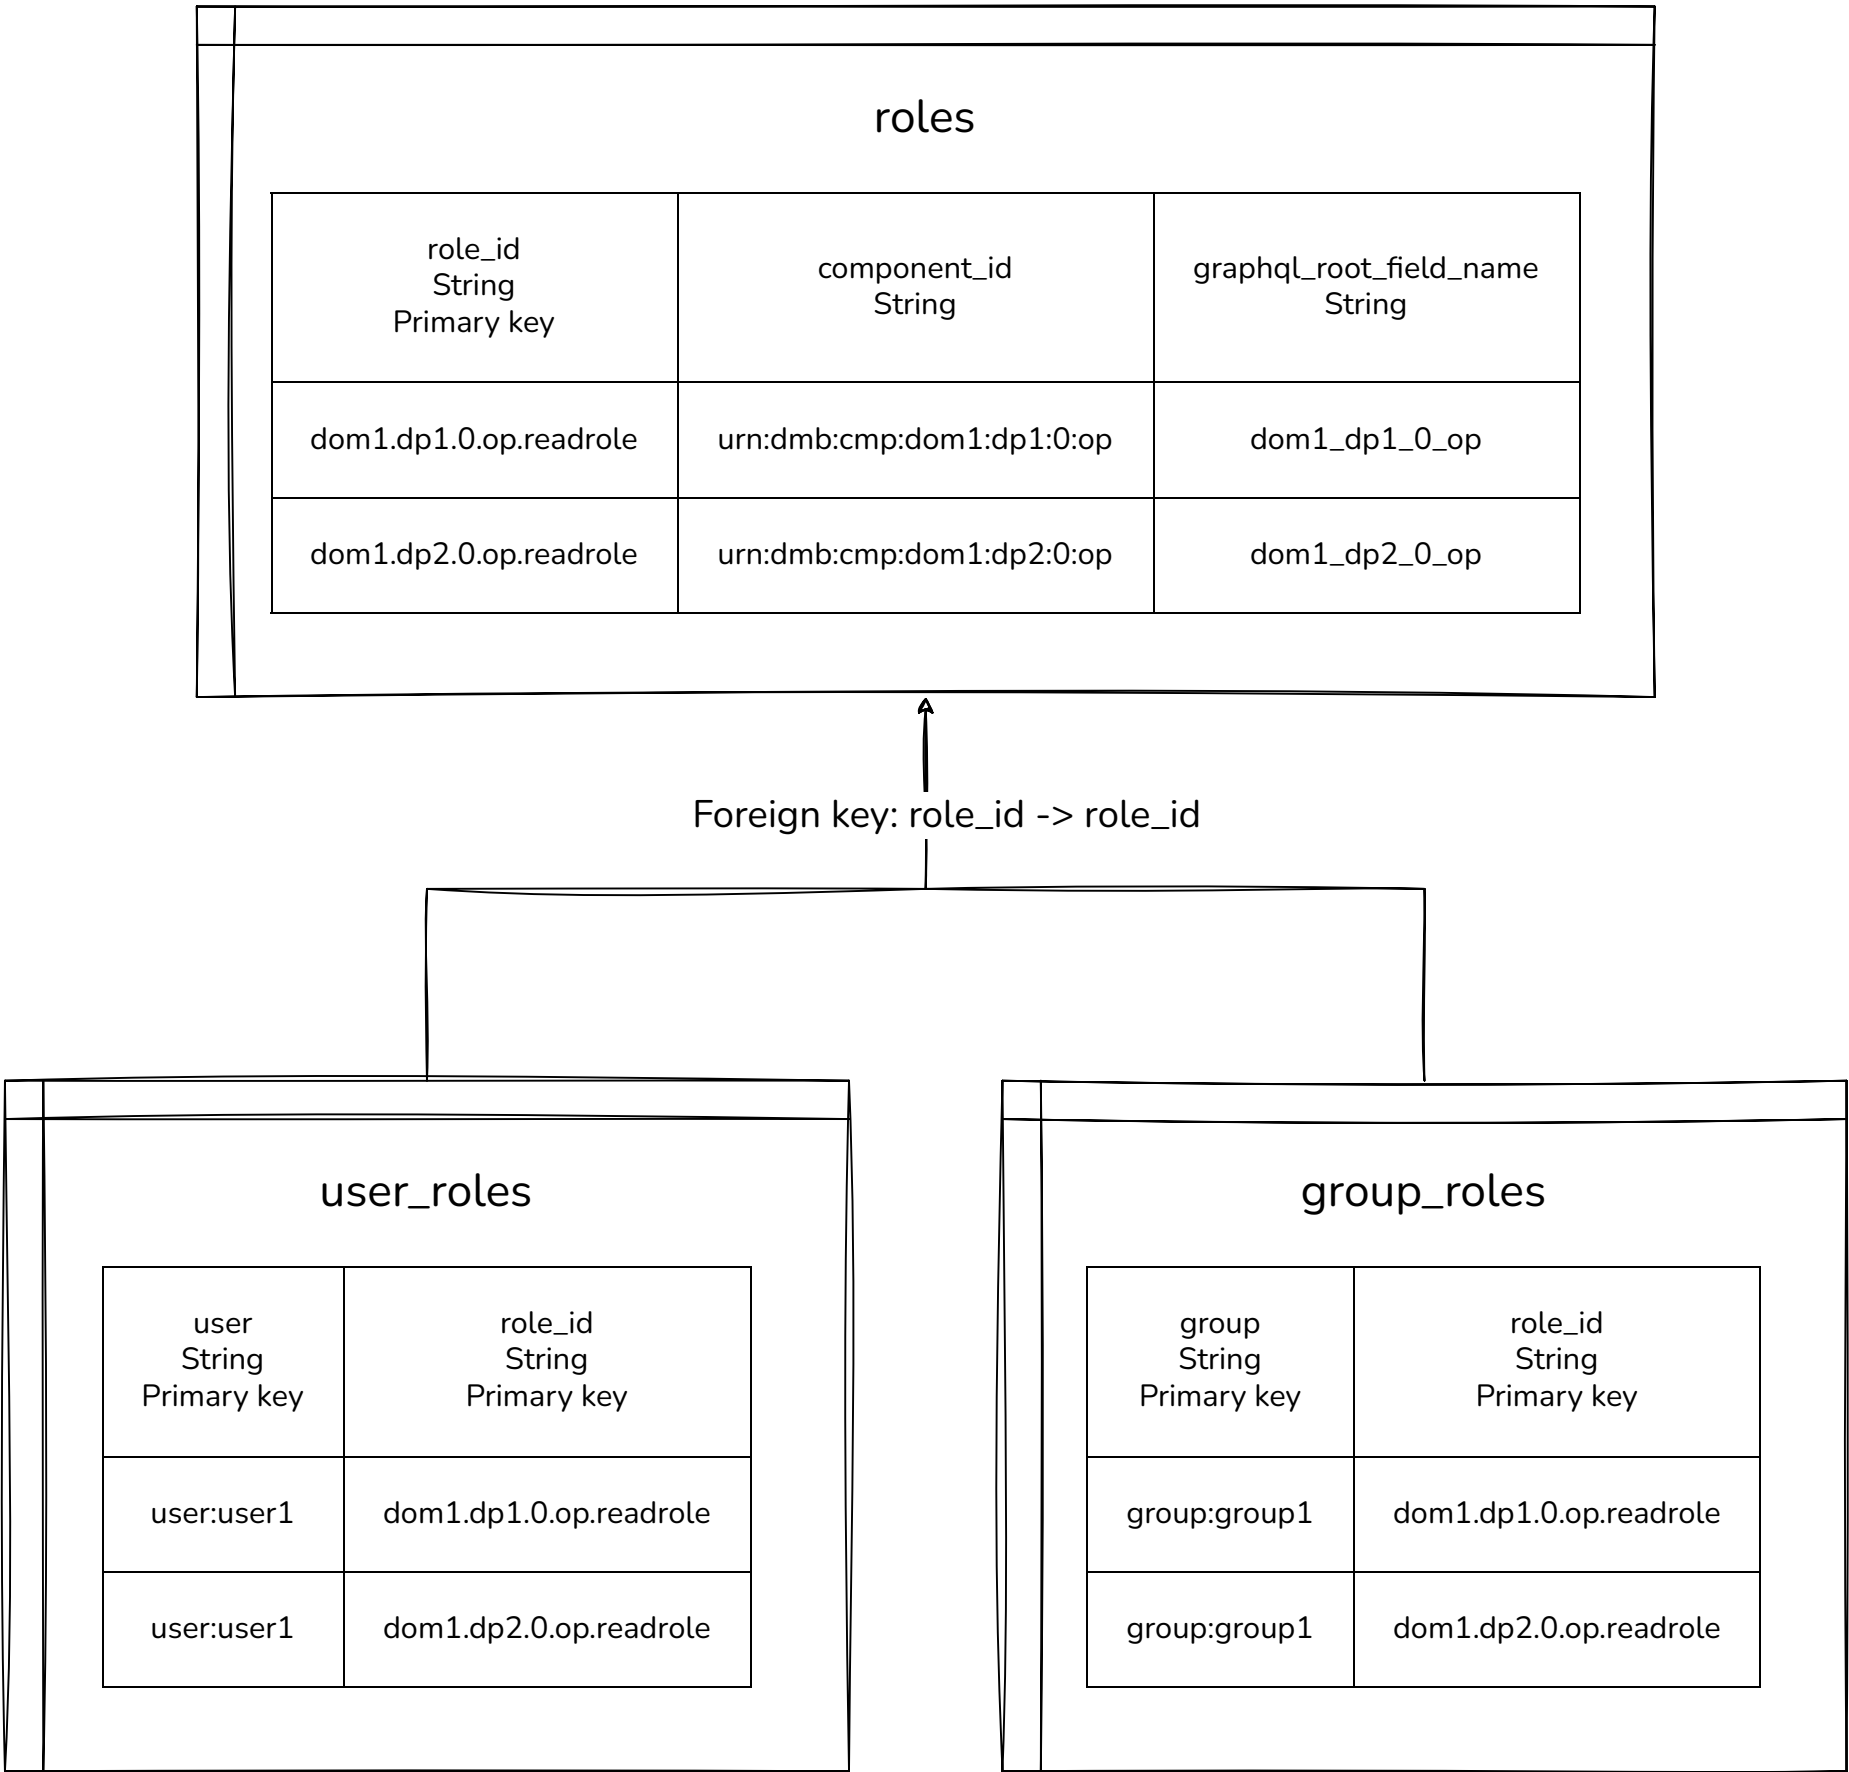
The roles table contains the role id and other information on each role, namely the corresponding component id and GraphQL root field; this is used in sanity checks to ensure multiple components are not using the same role, as well as for retrieving the correct role.

The user\_roles & group\_roles types map user & groups to sets of roles. The roles for a user are the set union of the user's roles (ie assigned to the user) and group roles (ie assigned to any of the user's groups).

The user and group identifiers in these tables are in the format used by Witboost (ie, user:user1 & group:group1 ); the translation from the format used by the IdP (and thus used in the JWT) to this format has to be done by the Webhook.

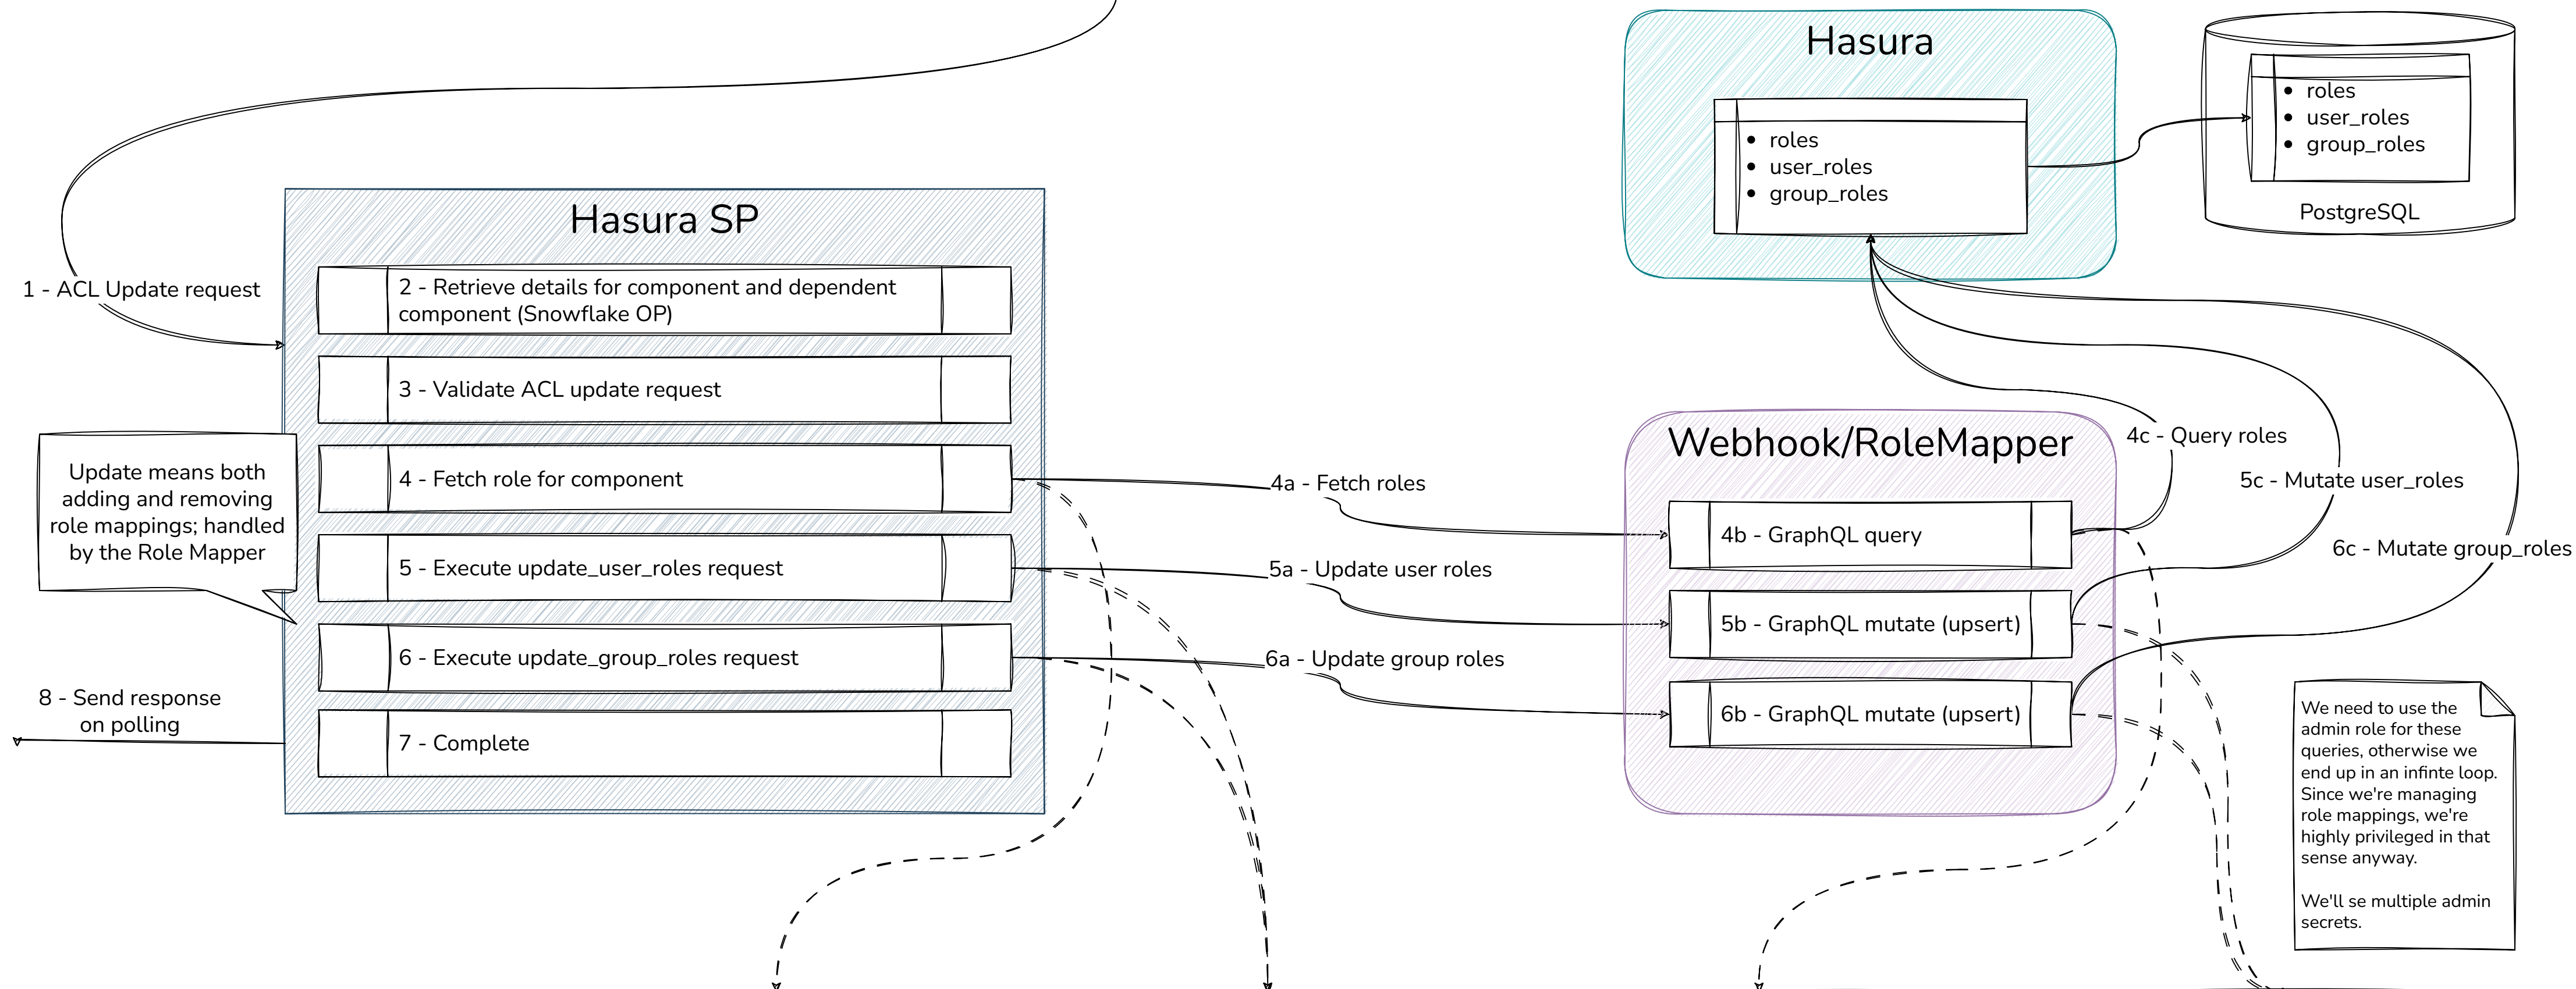

Data is stored in tables in underlying PG database. Permission to read/write is needed only by the role mapping part of the microservice.

These tables are tracked in Hasura and queried/manipulated via their GraphQL endpoints.



Mapping queries (which query by user/group) are far more frequent than updates (which update by role id), hence the decision to have the user/group as the first PK column, as it makes lookups slightly faster.





**Validations:**

- Check data source exists on Hasura
- Check Snowflake OP is from the same DP
- Check configuration for root fields is valid for GraphQL

```
GET
v1/role/component_id/urn:dmb:cmp:dom1:dp1:0:op

The provisioner can actually skip this if he knows the role name (eg because it is based on a naming convention).
```

```
For users:
PUT v1/user_roles
Content-Type: application/json
{
  "role": "dom1.dp1.0.op.readrole",
  "users": ["user:user1", "user:user2"]
}

For groups:
PUT v1/group_roles
Content-Type: application/json
{
  "role": "dom1.dp1.0.op.readrole",
  "groups": ["group:group1", "group:group2"]
}
```

```
POST v1/graphql
Content-Type: application/json
X-Hasura-Role: admin
X-Hasura-Admin-Secret: myadminsecretkey
{
  "query": "query GetRole($component_id: String!) {\n  roles(where: {component_id: {_eq: $component_id}}) {\n    component_id\n    graphql_root_field_name\n    role_id\n  }\n}",
  "variables": {
    "component_id": "urn:dmb:cmp:dom1:dp1:0:op"
  },
  "operationName": "GetRole"
}
```

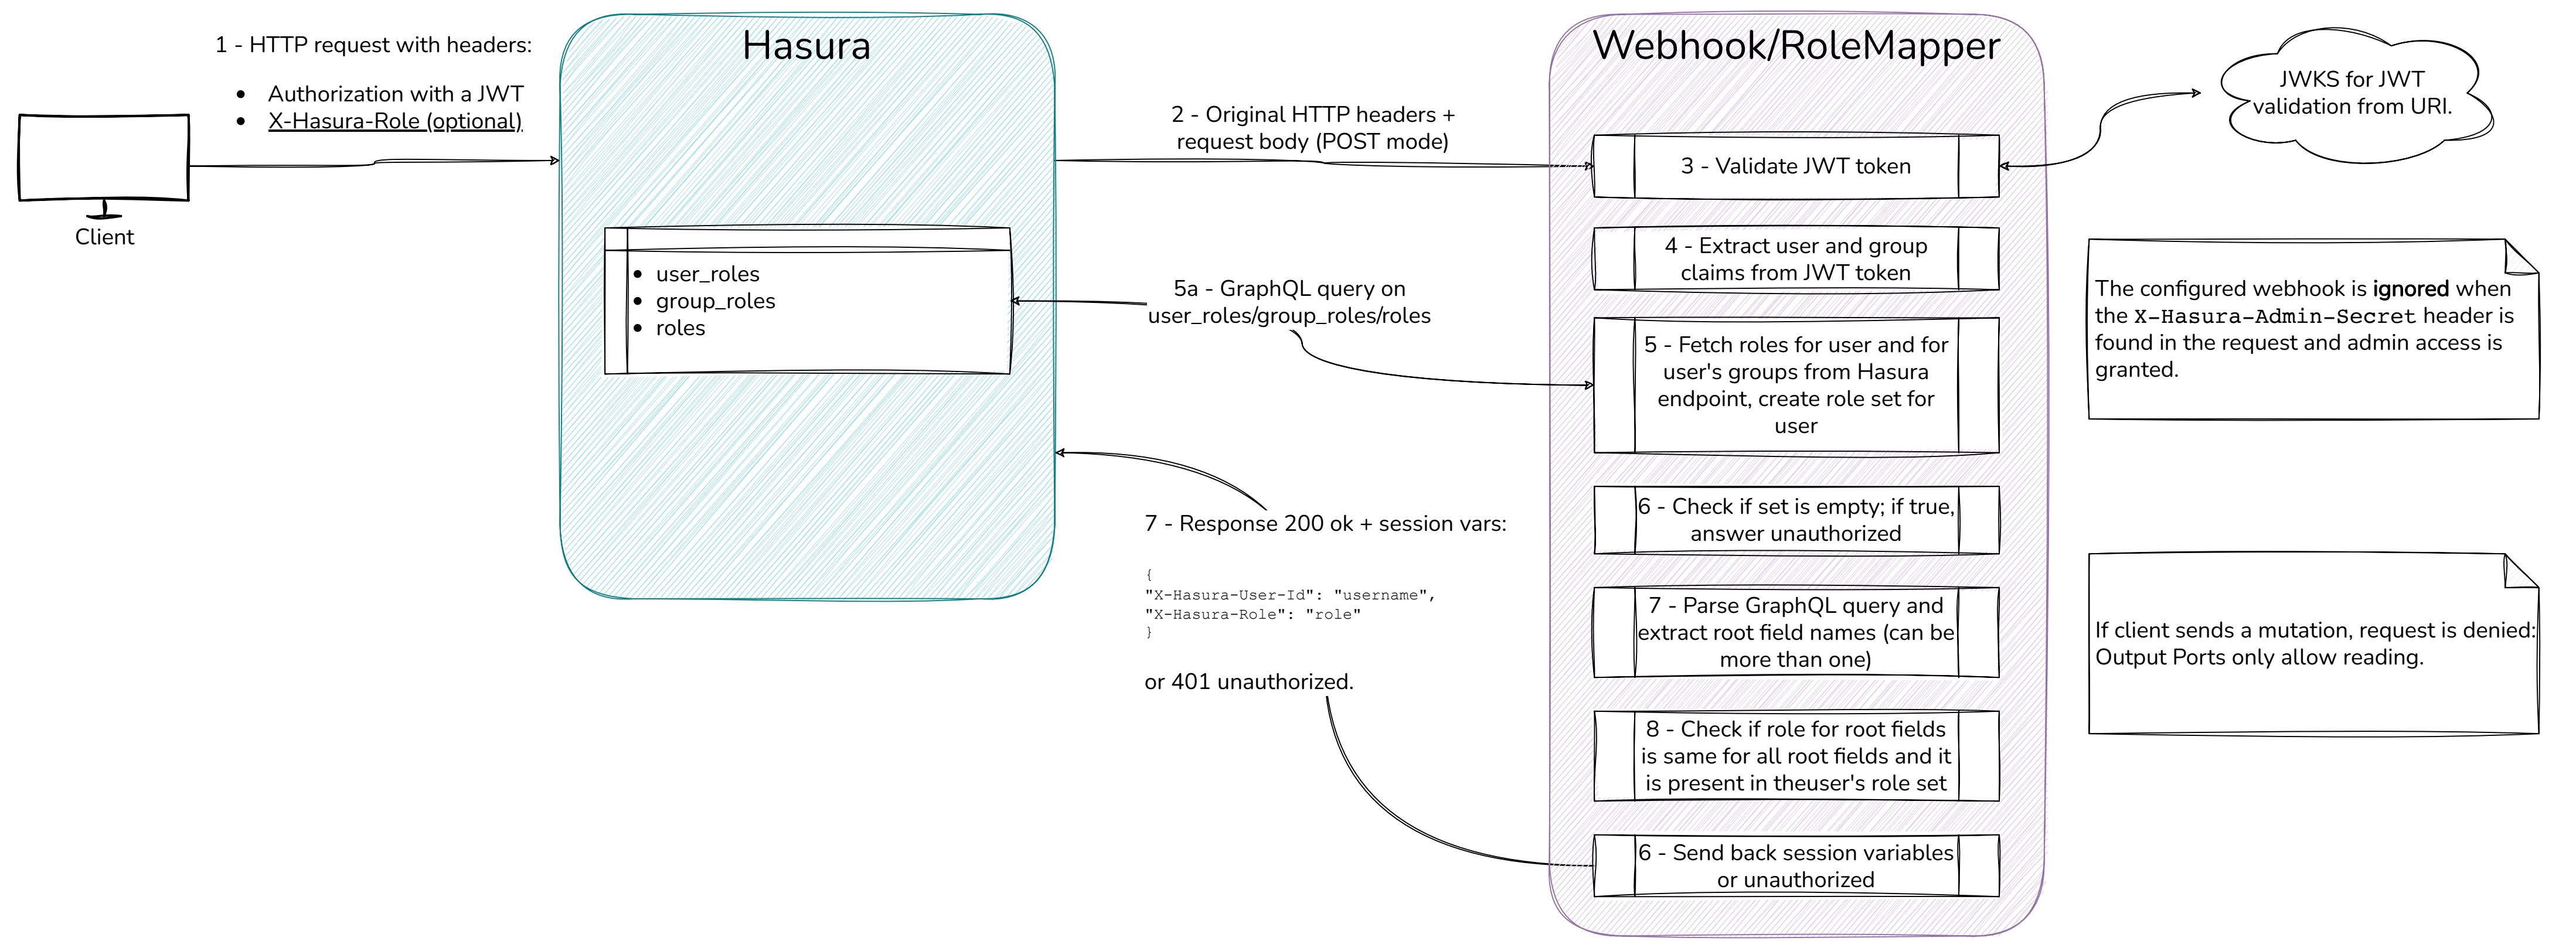
```
POST v1/graphql
Content-Type: application/json
X-Hasura-Role: admin
X-Hasura-Admin-Secret: myadminsecretkey
{
  "query": "mutation UpdateUserRole($user: String! $role_id: String! $last_update: timestampz = now!) {\n  insert_user_roles(objects: {last_update: $last_update, role_id: $role_id, user: $user}, on_conflict: {constraint: user_roles_pkey, update_columns: last_update, where: {}}) {\n    returning {\n      last_update\n      role_id\n      user\n    }\n  }\n}",
  "variables": {
    "role": "dom1.dp1.0.op.readrole",
    "user": "user:user1"
  },
  "operationName": "UpdateUserRole"
}
```

Above mutation is repeated for every user (can use bulk queries); group updates work the same way, only on group\_roles.

If a user & group are not mapped to a role anymore according to the request, we must also remove the previous mapping (omitted to avoid excessive clutter)



We use webhook mode as JWT mode requires customization on the IdP side to add Hasura's custom claims, and with some IdP applying custom logic to generate these claims is hard to do; also it would be too technology specific.



References:

- <https://hasura.io/docs/latest/auth/authentication/webhook/>
- <https://github.com/hasura/graphql-engine/tree/master/community/boilerplates/auth-webhooks>

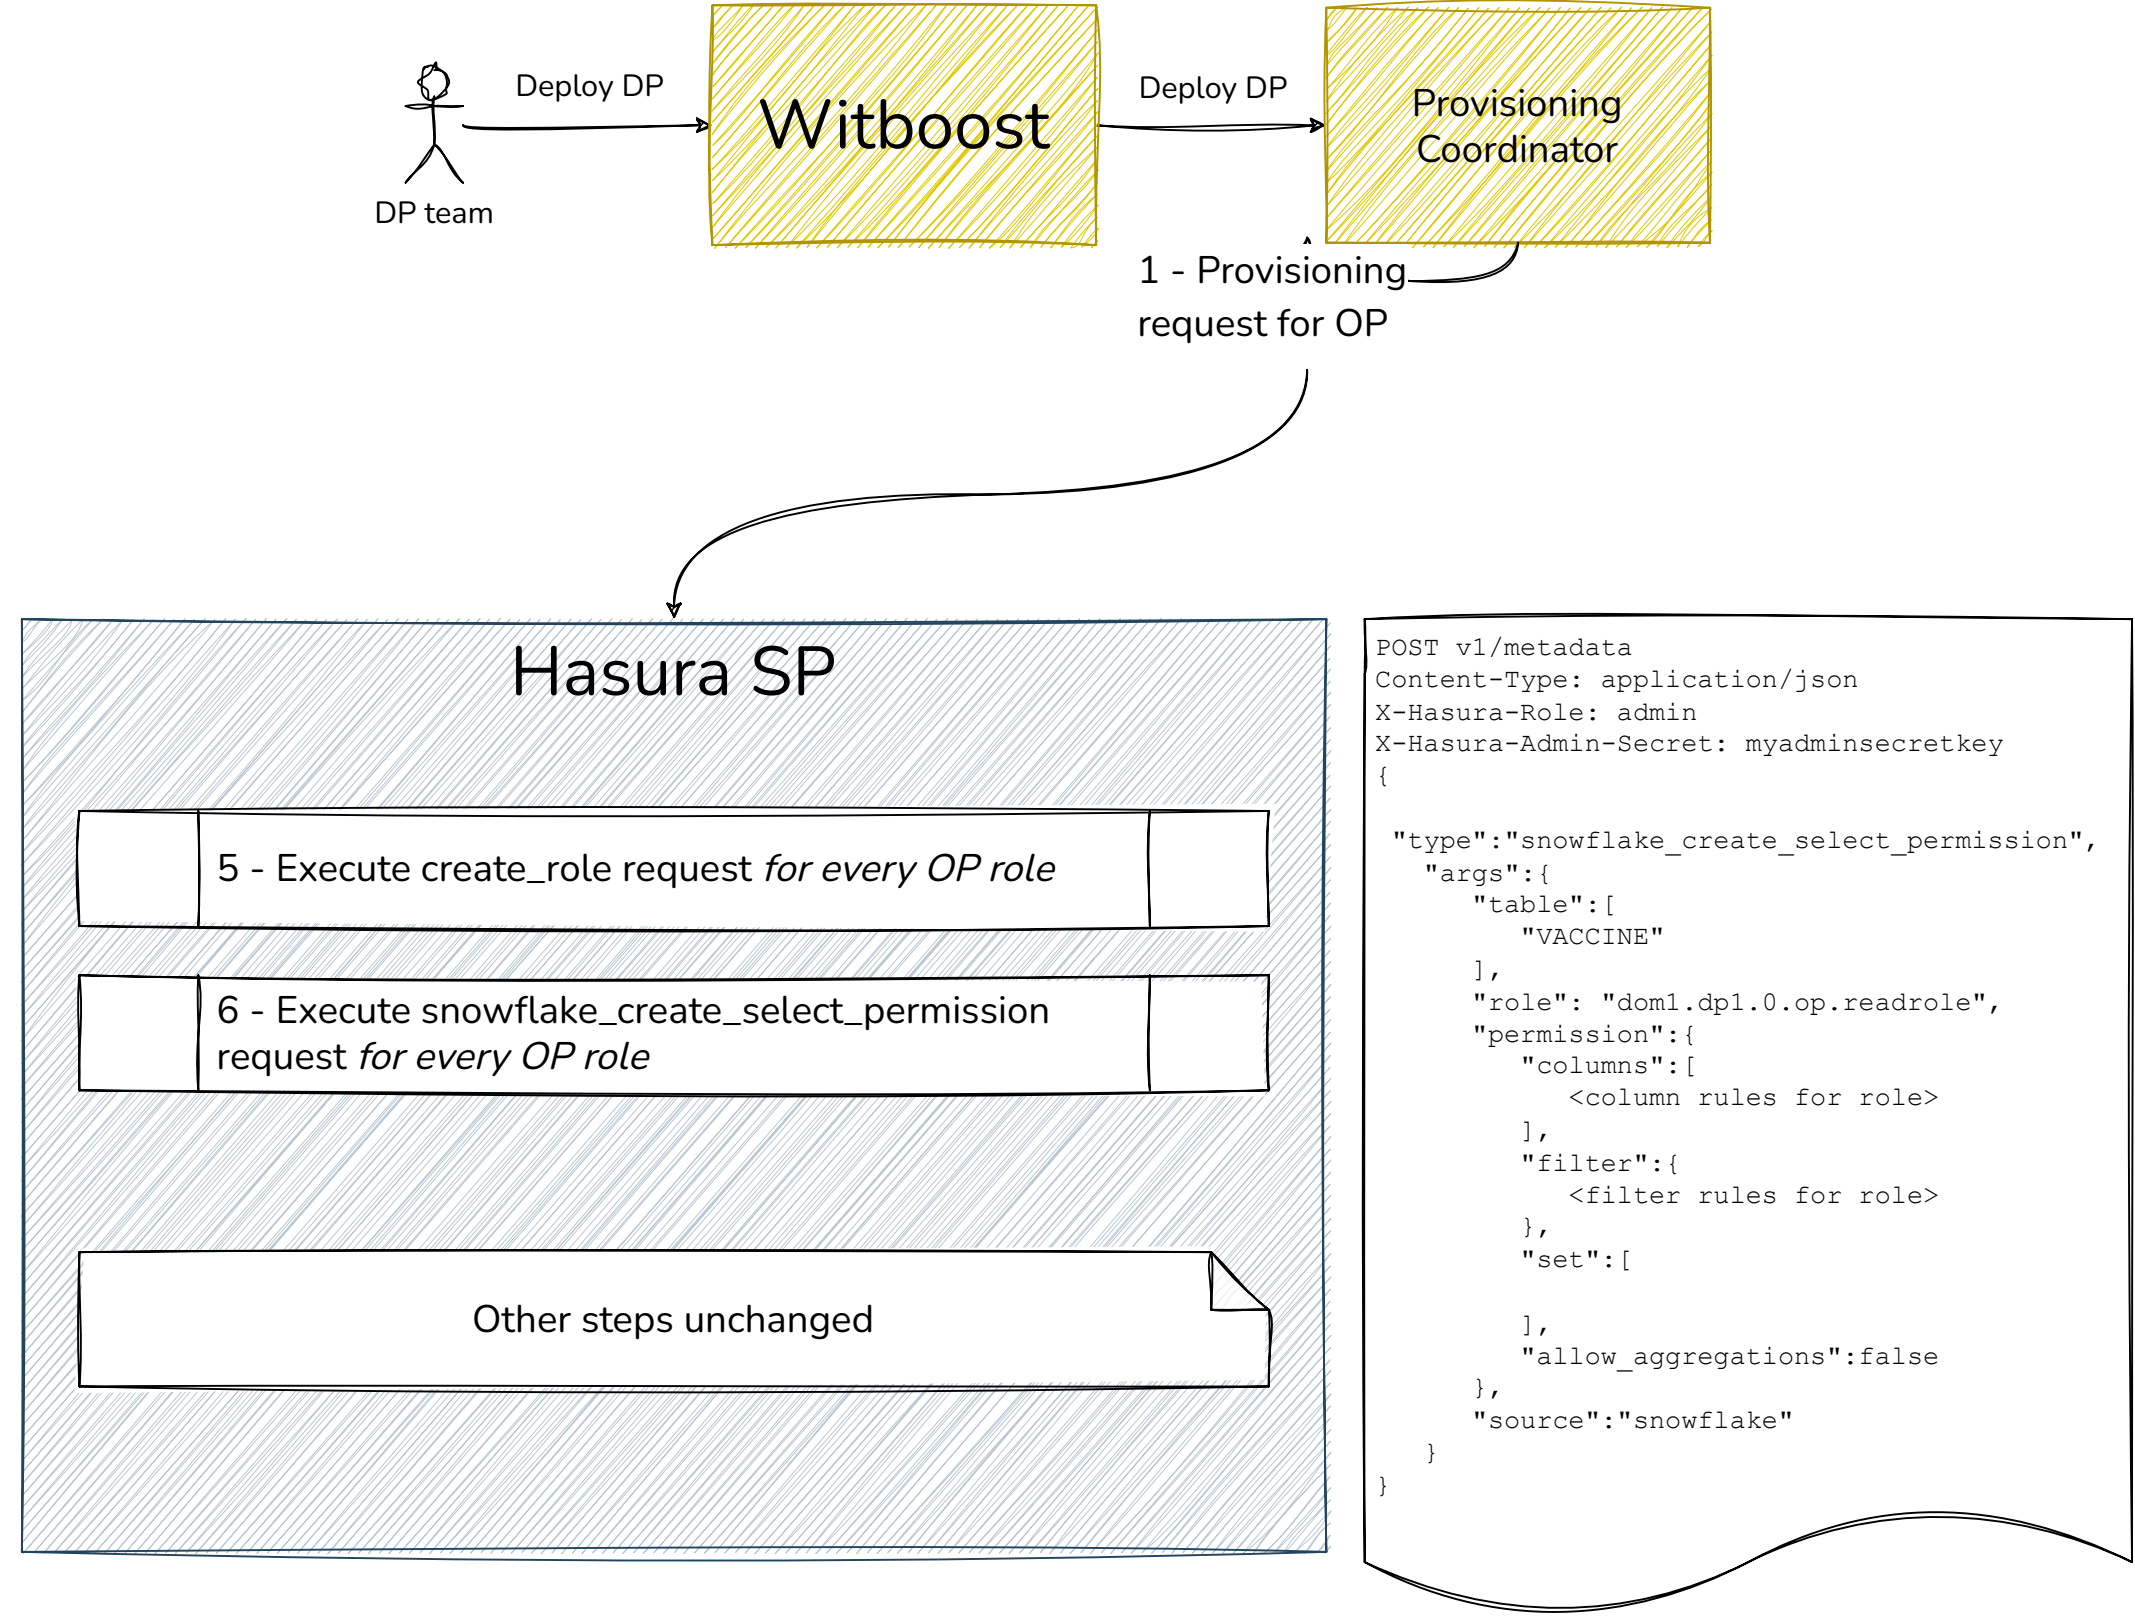


Out of scope for first release

Permissions discussed up to now are based on 1-1 mapping between Output Port and a role. This enables control of data access in an all-or-nothing fashion.

To enable more granular control, multiple roles can be used.

Provisioning (simplified)



ACL update (simplified)

